

# An Simple Cache System With Akka Cluster

For Scala Meetup in Shenzhen 2018.09.02

---

凤凰木 <weiwen@weiwen.org>

2018.09.02

# Content

业务需求

架构设计

实现

数据结构

一个基本的例子

分发集群化

一个分发的例子

组合的例子

问题与总结

## 业务需求

---

## 业务场景 1

- 一堆复杂的配置信息，存在 DB 里
- 配置不定期更新
- 一组应用服务器，需要实时读取这些配置信息

## 业务场景 2

- 若干张基础数据表，存在 DB 里
- 数据不定期更新
- 一组应用服务器，需要实时读取这些基础数据

## 业务场景 3

- 若干核心业务逻辑的报表，存在 DB 里
- 报表不断更新
- 一组应用服务器，需要实时读取这些报表
- 从 DB 读取报表的代价比较大，比如需要聚合

## 业务需求


1. 一组高性能的服务,  $2 \leq instances \leq 100$
2. 高速读取来自各种数据源的数据
3. 读取压力可能大到不允许有网络 IO
4. 数据不断被更新
5. 读取的代价可能比较大, 尤忌并发读取
6. 数据不要求强一致性
7. 数据量不是特别巨大, 远小于 JVM Heap

# 架构设计

---



## 最早的做法

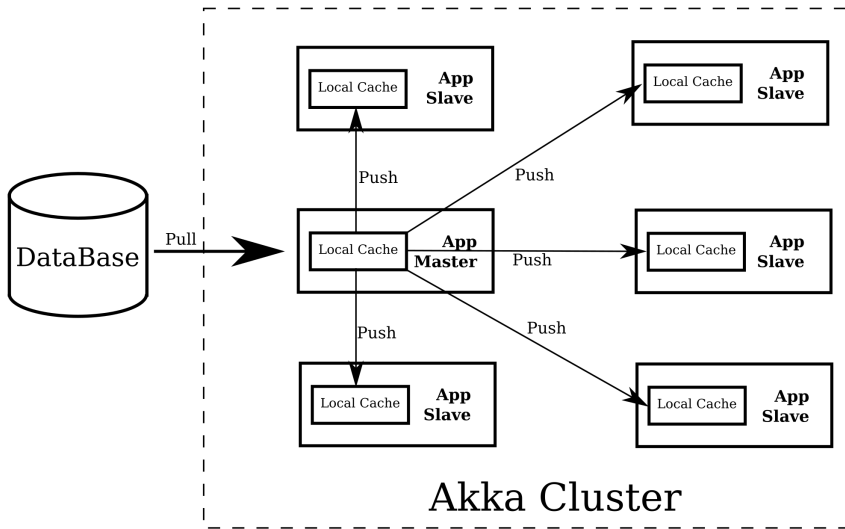


# 最早的做法

弊端：

- 数据源压力线性增长

## 改进的方案



## 各组件要点

- 缓存作为一个子系统，嵌入应用内
- 使用 Akka Cluster 来组成集群
- 使用 Akka Cluster Singleton Actor 作为 Master Node
- Akka Distributed Publish Subscribe 实现缓存的获取和分发逻辑
- Protobuf 作为数据传输格式

实现

---

# Data Structure: Protobuf

```
file src/main/protobuf/cache_meta.proto:
```

```
1  syntax = "proto3";
2
3  package simple.idl.cache;
4
5  enum CacheCate {
6      EmptyT    = 0;
7      SimpleT   = 1;
8      ComplexT  = 2;
9  }
10
11 message PackedCachePair {
12     bytes key    = 1;
13     bytes value  = 2;
14 }
15
16 message PackedCache {
17     int64 version = 1;
18     repeated PackedCachePair cc = 2;
19 }
```

# Data Structure: Protobuf, Why?

Questions:

1. Why **Protobuf** ?
2. Why **PackedCachePair** and **PackedCache** ?

## Data Structure: concurrent.TrieMap

```
1 scala.collection.concurrent.TrieMap
```

- A concurrent hash-trie or TrieMap is a concurrent thread-safe lock-free implementation of a hash array mapped trie.
- It has particularly scalable concurrent insert and remove operations and is memory-efficient.
- It supports  $O(1)$ , atomic, lock-free snapshots which are used to implement linearizable lock-free size, iterator and clear operations.
- The cost of evaluating the (lazy) snapshot is distributed across subsequent updates, thus making snapshot evaluation horizontally scalable.



## Data Structure: TrieMapCache

```
1 import scala.collection.concurrent.TrieMap
2 import simple.idl.cache.CacheCate
3
4 trait TrieMapCache[K, V] {
5
6     def underlying: TrieMap[K, V]
7
8     def cate: CacheCate
9
10    var version: Long = 0L
11    // var version: java.util.concurrent.atomic.AtomicLong = ???
12
13 }
```

## Data Structure: TrieMapCacheOps, 1

```
1  trait TrieMapCacheOps[K, V] {  
2  
3      this: TrieMapCache[K, V] =>  
4  
5      def setVersion(v: Long): Unit = {  
6          this.version = v  
7      }  
8  
9      def incVersion(): Long = {  
10         this.version += 1L  
11         this.version  
12     }  
13  
14 }
```

## Data Structure: TrieMapCacheOps, 2

```
1  trait TrieMapCacheOps[K, V] {  
2  
3      this: TrieMapCache[K, V] =>  
4  
5          // 按 key 获取  
6          def get(key: K, default: K => V): V = underlying.getOrElse(key, default(key))  
7  
8          // 按 key 获取  
9          def getIfPresent(key: K): Option[V] = underlying.get(key)  
10  
11         // 按 key 更新  
12         def put(key: K, value: V): Option[V] = underlying.put(key, value)  
13  
14         // 批量更新  
15         def putAll(m: Iterable[(K, V)]): TrieMap[K, V] = underlying ++= m  
16  
17         // 批量逐出, 按 keys  
18         def invalidateAll(keys: Iterable[K]): TrieMap[K, V] = underlying -= keys  
19  
20     }
```

## Data Structure: TrieMapCacheOps, 3

```
1  trait TrieMapCacheOps[K, V] {  
2  
3      this: TrieMapCache[K, V] ⇒  
4  
5          // 全量更新  
6          def updateAll(m: Iterable[(K, V)]): Unit = {  
7              val toDrops = underlying.keys.toVector.diff(m.map(_._1).toVector)  
8              invalidateAll(toDrops)  
9              putAll(m)  
10         }  
11  
12         def updateAllWithVersion(v: Long, m: Iterable[(K, V)]): Unit = {  
13             updateAll(m)  
14             setVersion(v)  
15         }  
16  
17         def updateAllAndIncVersion(m: Iterable[(K, V)]): Unit = {  
18             updateAll(m)  
19             this.incVersion()  
20         }  
21     }  
22 }
```

## Data Structure: InMemoryCache

```
1 trait InMemoryCache[K, V] extends  
2   TrieMapCache[K, V] with TrieMapCacheOps[K, V]
```

## Data Structure: PackableInMemoryCache, 1

```
1 import com.google.protobuf.ByteString
2 import simple.idl.cache.{PackedCachePair, PackedCache}
3
4 trait PackableInMemoryCache[K, V] extends InMemoryCache[K, V] {
5
6     val underlying: TrieMap[K, V] = TrieMap.empty[K, V]
7
8     def packKey(key: K): ByteString
9
10    def packValue(value: V): ByteString
11
12    def unPackKey(bytes: ByteString): K
13
14    def unPackValue(bytes: ByteString): V
15
16 }
```

## Data Structure: PackableInMemoryCache, 2

```
1 import simple.idl.cache.{PackedCachePair, PackedCache}
2
3 trait PackableInMemoryCache[K, V] extends InMemoryCache[K, V] {
4
5     def pack: PackedCache = {
6         val ps = underlying.toList.map { case (k, v) =>
7             PackedCachePair(packKey(k), packValue(v))
8         }
9         PackedCache(valid = 1, version = version, cc = ps)
10    }
11
12    def unPack(g: PackedCache): TrieMap[K, V] = {
13        val xs = g.cc.map { case PackedCachePair(k, v) =>
14            (unPackKey(k), unPackValue(v))
15        }
16        TrieMap(xs: _*)
17    }
18
19 }
```

## Data Structure: PbInMemoryCache

```
1 import com.google.protobuf.ByteString
2 import scalapb.GeneratedMessage
3
4 trait PbInMemoryCache[K <: GeneratedMessage, V <: GeneratedMessage]
5   extends PackableInMemoryCache[K, V] {
6
7   override def packKey(key: K): ByteString = key.toByteString
8
9   override def packValue(value: V): ByteString = value.toByteString
10
11 }
```



# SimpleCache: protobuf

file src/main/protobuf/simple\_cache.proto:

```
1  syntax = "proto3";
2
3  package simple.idl.cache;
4
5  message SimpleCacheKey {
6      int64 id      = 1;
7      int32 subId   = 2;
8  }
9
10 message SimpleCacheValue {
11     string name     = 1;
12     int64 impressions = 2;
13     int64 clicks    = 3;
14 }
```


## SimpleCache: scala, 1

```
1 import com.chinamobiad.adx.idl.cache.CacheCate
2 import com.google.protobuf.ByteString
3 import simple.idl.cache.{SimpleCacheKey, SimpleCacheValue}
4
5 object SimpleCache {
6
7     type K = SimpleCacheKey
8
9     type V = SimpleCacheValue
10
11 }
```

## SimpleCache: scala, 2

```
1 import com.chinamobiad.adx.idl.cache.CacheCate
2 import com.google.protobuf.ByteString
3 import simple.idl.cache.{SimpleCacheKey, SimpleCacheValue}
4
5 trait SimpleCache {
6
7     import SimpleCache._
8
9     val simpleCache: PbInMemoryCache[K, V] =
10
11         new PbInMemoryCache[SimpleCache.K, SimpleCache.V] {
12
13             val cate = CacheCate.SimpleT
14
15             def unPackKey(bytes: ByteString): K =
16                 SimpleCacheKey.parseFrom(bytes.toByteArray)
17
18             def unPackValue(bytes: ByteString): V =
19                 SimpleCacheValue.parseFrom(bytes.toByteArray)
20         }
21
22 }
```

# Master Actor Flow



# Configuration: build.sbt

file build.sbt:

```
1 libraryDependencies += Seq(  
2   "com.typesafe.akka" %% "akka-stream" % akkaV  
3   , "com.typesafe.akka" %% "akka-cluster" % akkaV  
4   , "com.typesafe.akka" %% "akka-cluster-tools" % akkaV  
5   , "com.twitter" %% "chill-akka" % "0.9.3"  
6 )
```

## Configuration: application.conf, 1

file src/main/resources/application.conf:

```
1 akka {  
2   actor {  
3     provider = cluster  
4     allow-java-serialization = off  
5     serializers {  
6       java = "akka.serialization.JavaSerializer"  
7       kryo = "com.twitter.chill.akka.AkkaSerializer"  
8       proto = "akka.remote.serialization.ProtoBufSerializer"  
9     }  
10    serialization-bindings {  
11      "scalapb.GeneratedMessage" = proto  
12      "scalapb.GeneratedEnum" = proto  
13      "java.io.Serializable" = kryo  
14    }  
15  }  
16 }
```

## Configuration: application.conf, 2

file src/main/resources/application.conf:

```
1 akka {  
2   cluster {  
3     seed-nodes = [  
4       "akka.tcp://simple@127.0.0.1:25510"  
5       , "akka.tcp://simple@127.0.0.1:25520"  
6       , "akka.tcp://simple@127.0.0.1:25530"  
7     ]  
8   }  
9   remote {  
10    enabled-transport = ["akka.remote.netty.tcp"]  
11    log-remote-lifecycle-events = off  
12    netty.tcp.port = 25510  
13  }  
14 }
```

## Reference: akka cluster config and serialization

### Akka Cluster:

- <https://doc.akka.io/docs/akka/current/cluster-usage.html>

### Serialization:

- <https://doc.akka.io/docs/akka/current/serialization.html?language=scala>



# PubSubCacheManagerT

```
1 import com.typesafe.scalalogging.Logger
2 import simple.idl.cache._
3
4 object PubSubCacheManager {
5
6     def masterNameOf(cate: CacheCate): String =
7         s"CacheMaster${cate.name}"
8
9     def topicNameOf(cate: CacheCate): String =
10         s"PubSubAcorTopic${cate.name}"
11
12 }
13
14 trait PubSubCacheManagerT[K, V] {
15
16     def cache: InMemoryCache[K, V]
17
18     protected val logger: Logger = Logger(getClass)
19
20 }
```

## PubSubCacheMaster, 1

```
1  trait PubSubCacheMaster[K, V] {  
2  
3      this: PubSubCacheManagerT[K, V] =>  
4  
5      import PubSubCacheManager._  
6  
7      // 每轮循环的间隔  
8      val tickInterval: FiniteDuration = 1000.milliseconds  
9  
10     // 把本地缓存 dump 成 protobuf 消息  
11     def dumpCache: scalapb.GeneratedMessage  
12  
13     // 从数据源获取数据, 更新本地缓存  
14     def updateMasterCache()(implicit ec: ExecutionContext): Future[_]  
15  
16 }
```

## PubSubCacheMaster, 2

```
1 import scala.concurrent.{ExecutionContext, Future}
2 import akka.actor._
3 import akka.stream._
4 import akka.stream.scaladsl._
5 import akka.cluster.pubsub.DistributedPubSub
6 import akka.cluster.pubsub.DistributedPubSubMediator._
7
8 trait PubSubCacheMaster[K, V] {
9
10   this: PubSubCacheManagerT[K, V] =>
11
12   import PubSubCacheManager._
13
14   def flow()(implicit system: ActorSystem
15               , ec: ExecutionContext): Flow[Unit, Any, NotUsed] = {
16     val publisher: ActorRef = DistributedPubSub(system).mediator
17     Flow[Unit]
18       .mapAsync(1)(_ => updateMasterCache())
19       .map(_ => cache.incVersion())
20       .map { _ =>
21         publisher ! Publish(topicNameOf(cache.cate), dumpCache)
22       }
23   }
```

## Reference: Distributed Publish Subscribe in Cluster

Distributed Publish Subscribe in Cluster:

- <https://doc.akka.io/docs/akka/current/distributed-pub-sub.html>

## PubSubCacheMaster, 3

```
1 trait PubSubCacheMaster[K, V] {
2   this: PubSubCacheManagerT[K, V] =>
3   def killSwitch()(implicit system: ActorSystem
4                     , mat: Materializer
5                     , ec: ExecutionContext): UniqueKillSwitch = {
6     RestartSource.withBackoff(minBackoff = 1.seconds
7                               , maxBackoff = 60.seconds, randomFactor = 0.2) { () =>
8       Source.fromFuture {
9         Source
10          .tick(0.seconds, tickInterval, ())
11          .via(flow)
12          .withAttributes(ActorAttributes.supervisionStrategy({
13            case NonFatal(e) =>
14              e.printStackTrace()
15              Supervision.Resume
16          })))
17          .runWith(Sink.ignore)
18      }
19  }
20  .viaMat(KillSwitches.single)(Keep.right)
21  .toMat(Sink.ignore)(Keep.left)
22  .run()
23 }
```

## Reference: RestartSource in Akka Stream

RestartSource:

- <https://doc.akka.io/docs/akka/current/stream/stream-error.html>

## PubSubCacheMaster, 4

```
1  trait PubSubCacheMaster[K, V] {  
2  
3      this: PubSubCacheManagerT[K, V] =>  
4  
5      private[this] class MasterActor extends Actor {  
6  
7          implicit val mat: ActorMaterializer = ActorMaterializer()  
8  
9          // 立即启动 Master 工作流  
10         val ks: UniqueKillSwitch =  
11             killSwitch()(context.system, mat, context.dispatcher)  
12  
13         override def postStop(): Unit = {  
14             ks.shutdown()  
15             mat.shutdown()  
16             super.postStop()  
17         }  
18  
19         def receive: Receive = PartialFunction.empty[Any, Unit]  
20  
21     }  
22 }
```

## PubSubCacheMaster, 5

```
1  trait PubSubCacheMaster[K, V] {  
2  
3      this: PubSubCacheManagerT[K, V] =>  
4  
5      def clusterSingleton()(implicit system: ActorSystem): ActorRef =  
6          system.actorOf(  
7              ClusterSingletonManager.props(  
8                  singletonProps = Props(new MasterActor)  
9                  , terminationMessage = 0  
10                 , settings = ClusterSingletonManagerSettings(system)  
11             )  
12             , name = masterNameOf(cache.cate)  
13         )  
14  
15     }
```



## Reference: Cluster Singleton

Cluster Singleton:

- <https://doc.akka.io/docs/akka/current/cluster-singleton.html>

# PubSubCacheClient

```
1  trait PubSubCacheClient[K, V] {  
2  
3      this: PubSubCacheManagerT[K, V] =>  
4  
5      import PubSubCacheManager._  
6  
7      def clientReceive: PartialFunction[Any, Unit]  
8  
9      def startClient()(implicit system: ActorSystem): ActorRef =  
10         system.actorOf(Props(new ClientActor))  
11  
12     class ClientActor extends Actor {  
13         override def receive: Receive = clientReceive  
14         override def preStart(): Unit = {  
15             val mediator = DistributedPubSub(context.system).mediator  
16             // 订阅消息  
17             mediator ! Subscribe(topicNameOf(cache.cate), self)  
18         }  
19     }  
20  
21 }
```

# SimplePbPubSubCacheManager

```
1  trait SimplePbPubSubCacheManager[K, V]
2      extends PubSubCacheManagerT[K, V]
3          with PubSubCacheMaster[K, V]
4          with PubSubCacheClient[K, V] {
5
6      override def cache: PackableInMemoryCache[K, V]
7
8      override def dumpCache: GeneratedMessage = cache.pack
9
10     def currentVersion: Long = cache.version
11
12     override def clientReceive: PartialFunction[Any, Unit] = {
13         // 需要检查版本号
14         case x@PackedCache(1, v, _) if v > cache.version =>
15             cache.updateAllWithVersion(v, cache.unPack(x))
16     }
17
18     def start()(implicit system: ActorSystem): Unit = {
19         startClient()
20         val _ = clusterSingleton()
21     }
22 }
23 }
```

# SimpleCacheManager

```
1 import scala.concurrent.{ExecutionContext, Future}
2 import simple.idl.cache.{SimpleCacheKey, SimpleCacheValue}
3
4 class SimpleCacheManager(
5     val cache: PbInMemoryCache[SimpleCacheKey, SimpleCacheValue]
6     ) extends
7     SimplePbPubSubCacheManager[SimpleCacheKey, SimpleCacheValue] {
8
9     val dao: SimpleCacheDAO = ???
10
11     override def updateMasterCache()(
12         implicit ec: ExecutionContext
13     ): Future[Unit] = {
14         val mF: Future[List[(SimpleCacheKey, SimpleCacheValue)]] =
15             dao.queryAll()
16         mF.map(cache.updateAllAndIncVersion)
17     }
18 }
```

# Cache

```
1 object Cache extends SimpleCache
2   with ComplexCache
3   with WxxCache
4   with XxxCache
5   with YyyCache
6   with ZzzCache
```

# SimpleService

```
1 object SimpleService {  
2  
3     // 如果找不到，就返回默认值 0L  
4     def findClicks(id: Long, subId: Int): Long =  
5         Cache.simpleCache  
6             .get(SimpleCacheKey(id, subId))  
7             .clicks  
8  
9 }
```

# CacheManager

```
1 object CacheManager {  
2  
3     def start()(implicit system: ActorSystem  
4                 , ec: ExecutionContext) = {  
5         new SimpleCacheManager(Cache.simpleCache).start()  
6         new ComplexCacheManager(Cache.complexCache).start()  
7         ...  
8     }  
9  
10 }
```

# Main

```
1 import akka.actor.ActorSystem
2 import akka.stream.ActorMaterializer
3 import scala.concurrent.ExecutionContext.Implicits.global
4
5 object Main extends App {
6
7     implicit val system: ActorSystem = ActorSystem()
8
9     CacheManager.start()
10
11 }
```



## 问题与总结

---

# Conclusion

- 弹性好, 可以很容易拓展
- 可靠性高, 即使集群脑裂也暂时无事
- 缓存没有过期机制, 但可以加
- 容量有限
- 总是全量更新, 有利有弊, 可以改进

完  
谢谢大家！